

::: {cover}

FluentEdge

Prototype Implementation Document

Subtitle: Sprint 4 - Local-first AWS-compatible MLOps prototype for AI-powered speaking feedback

Student / Sole CLC Team Member: Daniel Grijalva

Institution: Grand Canyon University

Course: SWE-452

Instructor: Ryan Woodward

Submission Date: July 20, 2026

Document Version: 4.0

Document Status: Final Sprint 4 submission document

Repository: [DanielAndi/fluentedge](#)

Sprint 4 Implementation Commit (Git SHA-1): 5c191a638110e28ceff15e259766bc1c04dc9d99

:::

Table of Contents

... {toc}

- 1. Project Overview
- 2. Requirements
- 3. Implementation Plan
- 4. Risk Assessment
- 5. Test and Validation Plan
- 6. Sprint 4 Project Documentation
- 7. References
- 8. Appendix A - Sprint 2 Requirements Baseline
- 9. Appendix B - Sprint 1 Risk Baseline
- 10. Appendix C - Current Evidence Index

...

1. Project Overview

1.1 Objective

FluentEdge demonstrates the first operational phase of an MLOps/cloud project for controlled English speaking feedback. A learner submits an expected phrase and a short audio recording. The service validates and temporarily stores the recording, loads an approved versioned model, and returns a score, a `pass` or `needs_review` label, confidence, feedback categories, model version, request ID, and latency.

The Sprint 4 objective is not to claim production-grade pronunciation assessment. It is to prove that the complete ML lifecycle can be executed and inspected consistently:

1. Validate and version source data.
2. Normalize audio and text deterministically.
3. Produce speaker-aware dataset partitions.
4. Build finite, versioned features.
5. Train and compare at least two candidate models.
6. Apply a measurable evaluation gate.
7. Register and explicitly approve a versioned model.
8. Deploy the approved model through a controlled API.
9. Monitor service health, latency, predictions, and dependencies.
10. Retain a rollback target and trace implementation evidence to requirements.

1.2 Measurable Prototype Goals

- Start the complete local stack from documented commands and confirm all required services are healthy.
- Validate a versioned dataset manifest and generate deterministic train, validation, and test assignments.
- Compare at least two baseline classifiers and block promotion below macro F1 `0.75`.
- Require human approval before a registered model can be treated as production-ready.
- Return a deterministic API response with safe errors, request IDs, and no raw audio in logs.
- Expose operational metrics through Prometheus and Grafana.
- Retain the previous approved model as a tested rollback target.
- Trace requirements to code, tests, commits, GitHub issues, reports, and screenshots.
- Record the full SHA-1 Git commit ID and SHA-256 digests for reproducibility-critical artifacts.

1.3 Prototype Scope

In scope	Out of scope
Controlled expected-phrase comparison using short audio clips	Unrestricted conversation or a full language-learning platform
WAV, FLAC, MP3, and M4A input when decodable	Unsupported or undecodable media
10 MB and 30-second prototype input limits	Long-form audio processing
Data schema, validation, normalization, and speaker-aware splitting	Claims that synthetic fixtures represent real student populations
Acoustic and text-derived features	Clinical or diagnostic pronunciation assessment
Logistic regression and random forest baselines	Large neural models unless added as a separately labeled experiment
MLflow experiment tracking, registration, and approval	Automatic model promotion without human review
FastAPI learner interface and OpenAPI contract	Enterprise authentication or a commercial mobile application

In scope	Out of scope
Docker Compose, LocalStack S3, MLflow, Prometheus, and Grafana	Always-on paid AWS infrastructure
AWS SAM health adapter and AWS-compatible target mapping	A claim that managed AWS deployment has occurred without evidence
GitHub Issues, Project #4, Actions, and stored evidence	Jira and Confluence as active systems; GitHub is the documented replacement

1.4 Stakeholders, Roles, and Responsibilities

Stakeholder	Role	Responsibilities
Daniel Grijalva	Project owner and sole CLC team member	Planning, architecture, implementation, model development, testing, documentation, GitHub governance, evidence, and demonstration
Ryan Woodward	Instructor and evaluator	Review technical rigor, evidence, completeness, risk treatment, reproducibility, and measurable success criteria
Learner / demonstration user	Primary user	Provide an expected phrase and short audio attempt; interpret plain-language prototype feedback
Future ML engineer	Technical maintainer	Curate licensed real data, reproduce training, evaluate bias and drift, approve models, and maintain model cards
Future DevOps operator	Platform maintainer	Operate containers, storage, observability, secrets, rollback, CI/CD, and an optional approved AWS deployment

CLC status: This is a one-person CLC. Every implementation, testing, documentation, evidence, and demonstration responsibility is assigned to Daniel Grijalva; there are no omitted team-member contributions.

2. Requirements

Sprint 2 reference: Appendix A reproduces the 65 functional and non-functional requirements from `FluentEdge_Sprint2_Design_and_Requirements_Specification.md`. No Sprint 2 identifier is deleted or renumbered. The change log below records implementation clarifications, measured gaps, and Sprint 4 documentation updates.

2.1 Requirement Change Log

Change ID	Area	Updated interpretation	Effect on implementation or validation
RC-01	Deployment profile	The local Docker implementation is authoritative for the class proof of concept. AWS remains a target mapping.	No managed AWS deployment is claimed without cloud evidence.
RC-02	Project tracking	GitHub Issues and GitHub Project #4 replace Jira and Confluence for task and evidence management.	FR-PM identifiers remain unchanged; evidence uses GitHub links and fields. A waiver should be attached if the instructor requires the named tools.
RC-03	Dataset	The implementation uses 48 synthetic fixtures only.	Functional tests may pass, but real-world validity and subgroup claims remain blocked.
RC-04	Model claims	Macro F1 <code>1.0</code> is synthetic functional evidence.	The result cannot be presented as pronunciation accuracy or generalization.
RC-05	Performance	Model-only latency and end-to-end API latency are separate measurements.	A fresh Python 3.11 run measured 30 warm <code>/predict</code> requests at p95 <code>84.749 ms</code> , below the <code>2,000 ms</code> local target.
RC-06	Partitioning	Six speaker groups previously produced an empty validation partition under floor-based splitting.	The allocator now guarantees every requested partition is nonempty when enough speaker groups exist; six groups produce 4/1/1 speaker groups.
RC-07	Dependencies	The audit reports MLflow/pyarrow advisories.	Security acceptance requires upgrade, remediation, or documented risk acceptance.
RC-08	Governance	Evidence documents conflict with open issue states and unchecked acceptance boxes.	Issue status, completion notes, and evidence fields must be reconciled before final acceptance.
RC-09	Cryptographic hashes	Git object IDs use SHA-1; dataset, model, report, and evidence checksums use SHA-256.	Every reproducibility record names the algorithm, digest length, and authoritative value.
RC-10	Latency target	The authoritative local target is warm <code>/predict</code> p95 at or below 2.0 seconds across at least 30 requests.	The earlier 800 ms AWS target is retained only as an aspirational future optimization.
RC-11	Repository freshness	Repository <code>main</code> was at <code>4e46d3d3047ad0ebaa05caec05de343d12ce3d8a</code> before the Sprint 4 hardening work.	Sprint 4 implementation commit <code>5c191a638110e28ceff15e259766bc1c04dc9d99</code> and documentation fix <code>e68e8f0d766701f76747b221ebd70f2ffeb2ab62</code> are pushed and visible.
RC-12	Evidence freshness	Executable and local-service screenshot evidence was refreshed on July 20, 2026 local time.	E-11 stores validation/test evidence; E-06 through E-08 contain fresh MLflow, Grafana, Prometheus, and learner-UI captures.

2.2 Reproducibility Requirement Clarification

Requirement	Sprint 4 clarification
FR-DI-006	The partition manifest shall record seed <code>42</code> , speaker/group assignments, and a SHA-256 digest of canonical manifest bytes.

Requirement	Sprint 4 clarification
FR-ML-002	<code>code_version</code> shall store the full 40-character Git SHA-1 commit ID; <code>dataset_manifest_sha256</code> shall store the 64-character SHA-256 digest.
FR-ML-007	The served artifact shall include a version and SHA-256 checksum; the API exposes the version while evidence retains the checksum.
FR-INF-006	Immutable artifact directories shall include the dataset version, run ID, and a SHA-256 inventory or digest file.
FR-PM-005	Requirement-to-commit traceability links the full Git SHA-1 commit ID, even when GitHub displays an abbreviation.
NFR-REL-002	Rollback evidence verifies the previous artifact SHA-256 checksum before loading it.

2.3 Verified Reproducibility Metadata

```
git_object_format: sha1
git_commit_sha1: 5c191a638110e28ceff15e259766bc1c04dc9d99
dataset_manifest_sha256: 81b47ea142995e7cf6361c8f7e2b0272c78ba66ee94dc8f8393fd56a3b33408e
model_artifact_sha256: c7016b8858bf0efa35c6f364cb2dc1307734d71c33975d299f8e8a0e5ae05e82
random_seed: 42
```


Integrity note: The production model digest was reverified by the Sprint 4 rollback exercise. The complete current inventory is stored in `docs/evidence/E-11-sprint4-validation/artifact-inventory.sha256` and must be regenerated after any listed file changes.

3. Implementation Plan

3.1 Technical Approach

Sprint 1 proposed an AWS-native architecture using Amazon S3, SageMaker Pipelines, SageMaker Model Registry, a FastAPI container on a SageMaker endpoint, SageMaker Model Monitor, CloudWatch, GitHub Actions, Jira, and Confluence. The implemented prototype preserves the same separation among data ingestion, training, registration, deployment, and monitoring while using a local-first, adapter-based architecture that is practical for the student budget and current course phase.

Application logic can use LocalStack S3 or a filesystem storage adapter, and the API receives environment-driven endpoints that can later map to AWS services. The training pipeline writes immutable run artifacts, records experiments in MLflow, requires explicit human approval, and serves only a versioned model. Prometheus scrapes application metrics, Grafana presents the operations dashboard, and GitHub Actions supplies repeatable quality, security, container, and documentation checks.



Sprint 1 Design-to-Prototype Alignment

Sprint 1 proposal choice	Implemented prototype choice	Alignment or documented change
Common Voice public-data subset	48 synthetic functional fixtures	Pipeline wiring is demonstrable, but public-data validity remains a blocking validation task.
Amazon S3	LocalStack S3 plus filesystem adapter	Preserves an S3-compatible contract without requiring paid cloud resources.
SageMaker Pipelines	Reproducible local pipeline and Make targets	Implements the same gated ingest-through-register lifecycle locally.
SageMaker Model Registry	MLflow tracking and model registry	Preserves version, metrics, artifact, and human-approval metadata.
FastAPI container on SageMaker endpoint	Local FastAPI Docker container	Preserves the REST inference contract and container portability; managed AWS deployment is not claimed.
SageMaker Model Monitor and CloudWatch	Prometheus and Grafana	Preserves operational monitoring for traffic, errors, latency, prediction distribution, confidence, and model version.
GitHub Actions or AWS CodeBuild	GitHub Actions	Preserves automated formatting, linting, testing, build, security, and documentation evidence.
Jira and Confluence	GitHub Issues, GitHub Project #4, and repository evidence	Implements the approved project-management replacement while keeping task, decision, risk, and evidence traceability.

Local-to-AWS Mapping

Local component	Current responsibility	Future AWS mapping
FastAPI container	Validation, inference orchestration, response contract	API Gateway plus Lambda or managed container service
LocalStack S3 / filesystem	Temporary upload and artifact storage	Amazon S3
Local pipeline	Validation, features, training, evaluation	SageMaker Processing, Training, and Pipelines

Local component	Current responsibility	Future AWS mapping
MLflow	Experiment tracking, registry, approval metadata	SageMaker Model Registry or approved managed registry
Prometheus / Grafana	Metrics and operations dashboard	CloudWatch metrics, logs, alarms, and dashboards
GitHub Actions	CI, security, container, and documentation checks	GitHub Actions with approved OIDC-based AWS access

3.2 Infrastructure Setup

1. Clone [DanielAndi/fluentedge](#).
2. Check out and verify full Sprint 4 implementation SHA-1 commit ID `5c191a638110e28ceff15e259766bc1c04dc9d99`.
3. Run `git rev-parse --show-object-format`; the expected result for this baseline is `sha1`.
4. Copy `.env.example` to `.env`, retain placeholder local credentials only, and never commit `.env`.
5. Use Python 3.11 or the documented Docker commands because the project requires Python `>=3.11, <3.12`.
6. Run `make setup` and then `make up` to build and start LocalStack, MLflow, FastAPI, Prometheus, and Grafana.
7. Run `make bootstrap` twice to prove idempotent S3-compatible bucket initialization.
8. Run `make health` and retain service-state output.
9. Generate fixtures, validate data, run the local pipeline, and approve the intended MLflow model version.
10. Set the approved model URI/version in `.env`, restart the API, and verify `/ready` and `/predict`.
11. Calculate SHA-256 digests for the dataset manifest, model artifact, model card, and evaluation report.
12. Execute the full test matrix and store the results before recording the demonstration.

Reproducible Checkout Commands

```
git clone https://github.com/DanielAndi/fluentedge.git
cd fluentedge
git checkout 5c191a638110e28ceff15e259766bc1c04dc9d99
git rev-parse --show-object-format
git rev-parse HEAD
```

Expected authoritative values:


```
sha1
5c191a638110e28ceff15e259766bc1c04dc9d99
```

3.3 Component Work Packages

Work package	Area	Execution	Output	Current state
WP-1	Data contract	Validate fields, duplicates, readability, duration, size, media type, and manifest version	Schema, validation report, negative fixtures	Complete on fixtures
WP-2	Reproducible preparation	Normalize audio/text, assign speaker-aware partitions, version features, calculate manifest SHA-256	Clean manifest, normalized audio, feature table	Nonempty deterministic split guard complete; public data remains open
WP-3	Model development	Train logistic regression and random forest with seed 42; compare metrics and latency	Evaluation report, joblib model, model card	Complete on fixtures
WP-4	Registry and approval	Log full Git SHA-1 commit ID, SHA-256 artifact digests, metrics, and approval state	MLflow run, approved version, checksums	Version 1 approved; hash metadata to standardize

Work package	Area	Execution	Output	Current state
WP-5	Inference API	Validate request, generate request ID, store temporarily, load approved model, predict, delete upload, return typed response	FastAPI routes and OpenAPI contract	Functional
WP-6	Learner interface	Provide upload/record controls and plain-language output	Static UI and screenshot	Functional baseline
WP-7	Observability	Export request, error, latency, label, confidence, version, and dependency metrics	Prometheus configuration and Grafana dashboard	Functional baseline
WP-8	CI and governance	Run format, lint, tests, secret scan, dependency scan, Compose validation, container smoke, and docs checks	Actions runs, issues, evidence index	CI, Container, and Docs green; audit has 32 findings; Project refresh blocked
WP-9	Validation hardening	Add licensed public data, valid partitioning, load/concurrency, failure, rollback, startup repetition, and dependency remediation	New reports and issue closure evidence	Split, load, concurrency, failure, rollback, and startup complete; real data and dependencies open
WP-10	Sprint 4 submission	Consolidate current status, requirements, risk, test plan, evidence, and video walkthrough	Markdown/PDF report and hosted video plan	Markdown/PDF and local evidence complete; hosted video remains owner action

3.4 Model Development and Promotion Workflow



Detailed steps:

1. Ingest a versioned manifest and reject blocking validation errors.
2. Normalize audio to mono 16 kHz PCM and normalize expected text deterministically.
3. Serialize the canonical manifest and calculate its SHA-256 digest.

- Assign speaker-aware partitions and abort if a mandatory partition is empty.
- Build finite, versioned features and preserve their schema.
- Train at least two candidates using recorded parameters and seed 42.
- Evaluate held-out speakers using macro F1, per-class precision/recall, confusion matrix, WER, CER, and latency.
- Block promotion below macro F1 0.75 or when any data-integrity gate fails.
- Log the full Git SHA-1 commit ID and SHA-256 dataset/model/report digests in MLflow and the model card.
- Require explicit human approval and retain the prior approved version for rollback.
- Deploy the approved version and validate end-to-end behavior before recording evidence.

3.5 Monitoring and Logging Plan

Concern	Signal	Cadence	Trigger	Required response
Availability	Dependency status, <code>/health</code> , <code>/ready</code>	Every scrape and demonstration checkpoint	Dependency down or approved model unavailable	Return controlled error; investigate by request ID; open issue after reproducible failure
Traffic	Request count and status	Continuous	Unexpected 4xx/5xx increase	Inspect error-code distribution and request IDs
Latency	End-to-end request latency	Continuous; p50/p95 review	Warm <code>/predict</code> p95 above 2.0 seconds	Profile decode, features, model loading, inference, and storage
Predictions	Label count and confidence distribution	Per request and sprint review	Material distribution shift	Freeze retraining; inspect data and model behavior
Model	Active version, approval tag, artifact SHA-256	Startup and every release	Unexpected version, missing approval, or checksum mismatch	Stop deployment and restore last verified approved version
Privacy	Metadata-only logs and retention cleanup	Per request and daily cleanup	Raw audio, token, transcript, or expired object found	Delete evidence, rotate credential if needed, open incident issue
Governance	CI state, open risks, requirement/evidence fields	Every pull request and sprint close	Failed gate or stale issue state	Block closure until evidence is reconciled

Logging rules:

- Log timestamp, request ID, route, status, duration, model version, and safe error code.
- Do not log raw audio, authorization headers, access tokens, or full transcripts.
- Record SHA-256 checksum failures as controlled integrity events without logging private artifact contents.
- Use structured JSON logs where supported.

3.6 Milestones, Deadlines, and Exit Evidence

Milestone	Target	Deliverable	Exit evidence	State
Sprint 1 - Proposal	June 28, 2026	Objective, scope, AWS-native technical direction, and eight initial risks	16-page proposal and baseline risk assessment	Historical
Sprint 2 - Design and local baseline	July 6-8, 2026	Requirements, architecture, code skeleton, local stack, CI, and evidence	Specification, reports, screenshots, Actions links	Implemented; issue states remain open
Sprint 3 - Prototype implementation	July 12, 2026	Implementation document, updated risk, test plan, and current prototype	18-page PDF and repository evidence	Completed baseline
Sprint 4 - Documentation and demo	July 20, 2026 submission	Updated implementation document, evidence audit, and video walkthrough	This Markdown/PDF plus uploaded video link	Local package complete; video and remote publication blocked

Milestone	Target	Deliverable	Exit evidence	State
Final - Release and presentation	August 9, 2026	Reproducible release, reconciled issues, final demonstration and evidence package	Tagged release, closed acceptance issues, final presentation	Planned

4. Risk Assessment

Appendix B reproduces the eight risks from the Sprint 1 proposal's page 13 risk assessment. The updated register below retains those concerns, adds risks discovered through implementation and instructor feedback, and converts the original qualitative ratings into a reviewable Probability x Impact scale. Scores 1-5 are Low, 6-10 are Moderate, 11-15 are High, and 16-25 are Critical. Daniel Grijalva owns every active response as the sole team member.

4.1 Updated Risk Register

ID	Risk	P	I	Score	Mitigation / contingency	Trigger
R-01	Synthetic-data validity	5	5	25 Critical	Use a licensed public speech subset with a completed dataset card; separate fixture metrics from validation metrics	Public-data evaluation cannot be reproduced or misses the quality gate
R-02	Overfitting or leakage	4	5	20 Critical	Require speaker-disjoint, non-empty partitions, canonical manifest SHA-256, and grouped cross-validation for small samples	Perfect fixture metrics, duplicate speakers, leakage, or empty validation partition
R-03	Accent or subgroup bias	4	5	20 Critical	Do not claim fairness; collect licensed metadata; report only groups with at least 100 examples; require human review	Material subgroup gap or insufficient metadata
R-04	End-to-end latency	2	4	8 Moderate	Keep the verified artifact cached and rerun the 30-request profile after model or feature changes	Warm /predict p95 exceeds 2.0 seconds; Sprint 4 result is 84.749 ms
R-05	Dependency vulnerabilities	4	4	16 Critical	Upgrade MLflow/pyarrow on a branch, rerun regression tests, and block final release on unresolved high-severity findings without acceptance	Dependency scan continues to report actionable vulnerabilities
R-06	Local/AWS parity	3	4	12 High	Keep storage adapters and environment configuration; test SAM path; avoid cloud claims without evidence	Behavior changes when endpoint or environment configuration changes
R-07	Rollback failure	2	4	8 Moderate	Reject a mismatched candidate, verify the approved SHA-256, atomically restore it, and load-test the restored artifact	Previous artifact cannot be verified or loaded after restoration
R-08	Privacy or retention failure	2	5	10 Moderate	Immediate post-inference deletion, 24-hour maximum cleanup, generated keys, log redaction, and retention tests	Raw audio remains, sensitive marker appears, or path traversal succeeds
R-09	Single-developer capacity	4	3	12 High	Prioritize blocking acceptance gates, automate evidence, time-box optional AWS/neural work, and maintain a reproducible runbook	Critical work slips or evidence and issue state diverge
R-10	Governance inconsistency	4	3	12 High	Update issue boxes, evidence fields, statuses, deadlines, and completion notes at each sprint close	Repository reports completion while issues remain open/evidence-required
R-11	Reproducibility hash ambiguity	3	4	12 High	Name SHA-1 for Git object IDs and SHA-256 for artifact integrity; store full digests and algorithm labels	Evidence records only "SHA," an abbreviated commit, or an unlabeled digest
R-12	Service availability	2	3	6 Moderate	Health checks, dependency timeouts, controlled 503 errors, persistent approved artifact, and two clean-start tests	A dependency blocks the API or startup is not repeatable
R-13	Cloud cost or credentials	2	3	6 Moderate	Keep required demo local, use placeholders, require manual approval and cleanup plan for any AWS test	Unexpected AWS charge, long-lived key, or automatic deployment
R-14	Model or data drift	3	4	12 High	Monitor prediction distribution, feature summaries, confidence, labels, and WER/CER when ground truth exists; open a retraining issue and retain rollback	Material input/prediction shift, quality degradation, or sustained confidence change
R-15	Evidence freshness / stale baseline	4	4	16 Critical	Rerun tests, health checks, performance, prediction, and screenshots immediately before recording; commit or clearly label the new evidence package	Submission relies entirely on July 8 results or claims work not present in the repository

4.2 Risk Changes Since Sprint 1

Direction	Risk area	Rationale
Mixed	Data validity and overfitting	The nonempty fixture split is corrected, but perfect synthetic scores still leave real-data generalization and leakage as critical risks
Reduced	Performance	The fresh 30-request run measured warm <code>/predict</code> p95 at 84.749 ms against the 2-second target
New	Dependency vulnerabilities	The stored audit reports advisories in pinned MLflow/pyarrow dependencies
New	Governance inconsistency	Sprint 2 issues remain open while newer evidence documents report completion
New	Hash ambiguity	Instructor feedback requires the exact SHA version; unlabeled or abbreviated hashes are no longer acceptable
Reduced	Evidence freshness	Python 3.11 test, coverage, performance, concurrency, startup, rollback, and audit evidence was refreshed on July 20
Reduced	Cloud cost and credentials	The required implementation remains local-first and has no automatic managed AWS deployment
Reduced	Privacy exposure	Generated storage keys, deletion, retention cleanup, path-safety tests, and log redaction exist
Unchanged	Bias, local/AWS parity, and drift	No real demographic dataset, managed AWS validation, or production traffic exists, so these risks remain unresolved

4.3 Risk Review Cadence

- Review the register before a sprint starts, before model approval, before recording the demonstration, and before closing a milestone.
- Open a GitHub issue for every triggered High or Critical risk and link affected requirement IDs.
- Do not lower a risk score merely because a control exists; lower it only after the control is tested and evidence is stored.
- Record accepted residual risk in the model card and release notes.
- Verify that every recorded digest identifies `sha1` or `sha256` and has the correct length.
- Rerun and timestamp executable evidence immediately before recording and submission.

5. Test and Validation Plan

5.1 Testing Objectives

- Verify that data ingestion rejects unsafe or invalid inputs and produces deterministic normalized artifacts.
- Verify that model training is reproducible, compares multiple candidates, and blocks promotion below the quality gate.
- Verify that the full Git SHA-1 commit ID and SHA-256 artifact digests reproduce the intended run inputs and outputs.
- Verify that only an explicitly approved, versioned model is served and that rollback restores the prior verified version.
- Verify the API contract, errors, privacy controls, storage cleanup, request IDs, and metrics.
- Measure end-to-end performance, concurrency, startup repeatability, and failure behavior rather than model inference time alone.
- Confirm that CI, evidence links, and GitHub issue states agree before PoC acceptance.
- Rerun core evidence immediately before the Sprint 4 demo so the recorded video reflects the submitted state.

5.2 Test Environment

Dimension	Configuration
Runtime	Python 3.11; the project excludes Python 3.12 or later for this baseline
Container platform	Docker Engine/Desktop with Compose and the named <code>fluentedge</code> network
Services	FastAPI :8000 , MLflow :5000 , LocalStack S3 :4566 , Prometheus :9090 , Grafana :3000
Test automation	pytest, pytest-asyncio, Ruff, pip-audit, Gitleaks, GitHub Actions, Compose validation, and container smoke tests
Model stack	NumPy, pandas, scikit-learn, librosa, soundfile, jiwer, joblib, and MLflow 2.18.0 baseline
Git identity	Implementation SHA-1 5c191a638110e28ceff15e259766bc1c04dc9d99 ; green follow-up e68e8f0d766701f76747b221ebd70f2ffeb2ab62
Artifact integrity	SHA-256 digest for manifest, features, model, model card, evaluation report, and evidence inventory
Evidence baseline	Executable reports refreshed July 20, 2026; authenticated service/browser screenshots remain from July 6-8

5.3 Test Data

Dataset	Location / source	Composition	Purpose
Synthetic manifest	tests/fixtures/synthetic/manifest.json	48 WAV clips; 6 synthetic speakers; 8 expected phrases; 24 <code>pass</code> and 24 <code>needs_review</code> labels	Unit, pipeline, API, and smoke tests only
Negative fixtures	Corrupt, too-long, oversized, unsupported, missing-prompt, and boundary inputs	Decode failure, duration, size, type, and request-boundary cases	Deterministic error-code validation
Current partition	Seed 42 and speaker-aware 70/15/15 implementation	4 train, 1 validation, and 1 test speaker group; 32/8/8 clips for the balanced fixture	Functional partition gate passes; public-data validity remains open
Planned public set	Licensed Mozilla Common Voice English subset or approved equivalent	Version, locale, license, checksum, date, clip count, duration, speaker distribution, and limitations recorded	Real-world validation and subgroup analysis
Operational load set	At least 30 warm valid requests plus invalid and concurrent batches	Representative 1-10 second clips and at least five simultaneous clients	p50/p95 latency, reliability, and corruption checks

Validation rule: Synthetic fixtures may satisfy functional pipeline tests, but they may not support claims of real-world speaking accuracy, demographic fairness, or production readiness.

5.4 Detailed Test Matrix

ID	Area	Test	Expected result	Method	Current evidence
T-01	Data	Valid manifest and WAV set	48 of 48 rows valid; no blocking issue	Automated	Pass in stored evidence
T-02	Data	Missing field, duplicate ID/path, unreadable audio	Each category identified and pipeline blocked	Automated	Implemented
T-03	Input	Unsupported, corrupt, above 10 MB, or above 30 seconds	Deterministic 4xx/413/415, safe message, request ID	Automated	Implemented; rerun required
T-04	Preprocessing	Audio normalization	Mono 16 kHz PCM with finite waveform	Automated	Implemented
T-05	Text	Whitespace, case, and punctuation cases	Exact golden normalized text	Automated	Implemented
T-06	Split	Same seed and speaker IDs	Identical assignment, no overlap, all mandatory partitions non-empty	Automated	Pass; six-speaker regression covered
T-07	Hashing	Repository identity	<code>git rev-parse --show-object-format</code> returns <code>sha1</code> ; full commit matches baseline	Automated/manual	Planned explicit evidence
T-08	Hashing	Manifest and artifact integrity	Recalculated SHA-256 values exactly match stored 64-character digests	Automated	Pass; inventory stored in E-11
T-09	Features	Build feature frame	Required schema exists; no NaN/Inf; version recorded	Automated	Implemented
T-10	Training	Train two candidates	Both fit from clean checkout using seed 42	Automated/integration	Pass on fixtures
T-11	Evaluation	Held-out candidate comparison	Macro F1, class metrics, confusion, WER/CER, and latency recorded	Automated/integration	Pass on fixtures
T-12	Gate	Candidate below macro F1 0.75	Promotion blocked and remediation decision recorded	Automated	Negative gate test planned
T-13	Registry	Unapproved versus approved version	Unapproved load fails; approved version is traceable	Automated/manual	Approval evidenced
T-14	API	<code>GET /health</code> and <code>/ready</code>	Schema valid; dependency/model states accurate; no secrets	Automated/smoke	Pass
T-15	API	Valid <code>POST /predict</code>	Required fields, label domain, request ID, version, latency, cleanup	Automated/smoke	Pass; latency gap
T-16	API	Missing prompt, bad media, decode failure	Stable safe error code, status, and request ID	Automated	Pass
T-17	Security	Filename path traversal	Generated object key; no escape from storage root	Automated	Pass
T-18	Privacy	Sensitive log markers	No raw audio, bearer token, authorization header, or full transcript	Automated	Pass
T-19	Retention	Expired and successful upload cleanup	Successful inference deletes immediately; fallback within 24 hours	Automated/manual	Baseline pass
T-20	Storage	Filesystem and LocalStack adapters	Put/get/list/delete/cleanup behavior equivalent	Automated/integration	Baseline pass

ID	Area	Test	Expected result	Method	Current evidence
T-21	Performance	30 warm <code>/predict</code> requests	p95 at or below 2.0 seconds for clips at or below 10 seconds	Load	Pass; p95 84.749 ms
T-22	Performance	30 <code>/health</code> requests	p95 below 250 ms	Load	Pass; p95 36.732 ms
T-23	Concurrency	Five simultaneous valid requests	No corruption, duplicate IDs, or uncontrolled errors	Load	Pass; five unique IDs and zero leftover uploads
T-24	Reliability	Disable storage or model/registry dependency	Controlled 503; API remains alive; failure metric increments	Failure injection	Pass; storage and model-load tests return controlled 503 contracts
T-25	Rollback	Reject bad candidate then restore approved artifact	Previous SHA-256-verified model is restored and loadable	Automated/integration	Pass; approved SHA-256 restored exactly
T-26	Startup	Two clean start/bootstrap/health cycles	API, storage, model, and dependency health pass twice	System	Pass; two fresh startup logs stored
T-27	Observability	Generate prediction activity	Dashboard shows traffic, latency, dependency health, and prediction label	Manual/system	Pass; Grafana, Prometheus, MLflow, and learner UI recaptured
T-28	CI	Push or pull request	Format, lint, unit, Compose, container, docs, and secret checks pass	Automated	Pass; CI, Container, and Docs workflows green at <code>e68e8f0</code>
T-29	Dependencies	Locked dependency audit	No unresolved Critical/High finding or signed risk acceptance	Automated/review	Fail; 32 findings in MLflow 2.18.0 and pyarrow 18.1.0
T-30	Governance	Trace requirement to issue, full commit ID, test, and evidence	Links resolve and issue completion state matches evidence	Manual/review	Reconciliation required
T-31	Usability	Learner completes one attempt	Plain-language result, non-clinical wording, understandable error recovery	Manual	Baseline UI reviewed
T-32	Freshness	Rerun submission evidence	Current timestamped test, health, predict, and dashboard evidence matches the demo	Manual/system	Local and GitHub Actions evidence refreshed; Project screenshot and video open

5.5 Current Evidence Summary

Evidence	Stored result	Interpretation
Unit test suite	45 passed, 2 skipped, 1 integration test deselected	Functional baseline and new hardening regressions pass on Python 3.11
Format and lint	57 files formatted; Ruff check passes	Local quality gates pass
Data validation	48 of 48 rows valid	Synthetic manifest is internally valid
Model comparison	Logistic regression and random forest both macro F1 1.0	Pipeline is functional; result is not real-world accuracy evidence
Model registry	Version 1 approved	Human approval path is present
Container smoke	Pass	API image builds and exposes health route
Service health	Six required checks pass	Local stack is demonstrable
<code>/health</code> latency	p95 36.732 ms across 30 samples	Passes the 250 ms gate
<code>/predict</code> latency	warm p95 84.749 ms across 30 samples	Passes the 2-second local gate for the fixture
Concurrency/startup/rollback	Five-way concurrency, two starts, and verified rollback pass	Reliability hardening is evidenced locally

Evidence	Stored result	Interpretation
Security	Prior Gitleaks evidence passes; fresh dependency audit finds 32 advisories	Dependency gate remains blocked pending upgrade or signed acceptance

5.6 PoC Success Criteria

Criterion	Measurable threshold	Gate	Required evidence
Functional pipeline	Clean manifest completes validate -> preprocess -> split -> features -> train -> evaluate -> register	Blocking	Pipeline result and immutable artifacts
Reproducibility	Full 40-character Git SHA-1 commit ID and matching SHA-256 digests reproduce run inputs/outputs using the recorded seed	Blocking	Run metadata, checksum inventory, rerun transcript
Model quality	At least two candidates; selected real-data model macro F1 at least 0.75 with no unresolved leakage or empty partition	Blocking	Evaluation report and model card
Approval and rollback	Only human-approved version serves; previous SHA-256-verified version restores successfully	Blocking	MLflow tags and rollback transcript
API correctness	Contract, negative, privacy, retention, storage, and path-traversal tests pass	Blocking	JUnit/coverage and API samples
Performance	Warm /predict p95 at or below 2.0 seconds across at least 30 requests; /health p95 below 250 ms; five concurrent requests without corruption	Blocking	Load-test JSON and dashboard
Reliability	Controlled dependency failure and two consecutive clean startups pass	Blocking	Failure-injection and startup logs
Security	Secret scan passes; no unresolved Critical/High dependency issue without documented acceptance	Blocking	Gitleaks and dependency audit artifacts
Observability	Required metrics are visible and associated with request/model version	Required	Prometheus/Grafana evidence
Reproducible operation	Reviewer can start, test, train, approve, and demonstrate on Python 3.11/Docker from the documented baseline	Blocking	Fresh-checkout run log
Governance	Requirements trace to current issue state, evidence, deadlines, and completion notes	Required	GitHub Project and audit evidence
Evidence freshness	Video and submitted report reference the same current baseline and timestamped reruns	Required	Fresh test logs, screenshots, and video link

5.7 Entry, Exit, and Defect Rules

- **Entry:** Code is on the recorded full SHA-1 commit; environment uses Python 3.11; fixtures and approved data are versioned; SHA-256 manifests exist; required services are healthy.
- **Exit:** Every Blocking criterion passes; evidence is stored; limitations are in the model card; related GitHub issues contain accurate completion notes.
- A failed Blocking test prevents PoC acceptance and cannot be converted to a pass through narrative explanation alone.
- A deferred optional requirement must identify owner, rationale, target sprint, and evidence that the deferral does not invalidate a mandatory criterion.
- Defects receive severity, affected requirement IDs, reproduction steps, expected/actual behavior, evidence, and retest result.
- **Freshness:** Any result shown in the video must be reproducible from the same repository baseline cited in the document.

Current validation decision: The repository demonstrates a functional prototype and now passes the local partition, API performance, health performance, concurrency, controlled-failure, startup, rollback, unit-test, lint, artifact-integrity, local observability, and GitHub Actions gates. Final PoC acceptance is still blocked by licensed real-data validity, dependency remediation or signed risk acceptance, authenticated GitHub Project reconciliation, and the hosted demonstration video.

6. Sprint 4 Project Documentation

6.1 Nature and Scope of Sprint 4 Tasks

Sprint 4 converts the existing functional prototype and Sprint 3 report into a current, reviewable submission package. The sprint scope includes documentation, evidence reconciliation, demonstration planning, and executable validation hardening. Remote GitHub state is reported separately because the configured `gh` token is invalid and no unauthenticated write is claimed.

Task area	Sprint 4 scope
Documentation consolidation	Update project overview, requirements, implementation plan, risk register, test plan, and appendices.
Repository audit	Verify current baseline commit, README setup path, issue status, CI/test evidence, and known limitations.
Rubric alignment	Explicitly document work scope, sole-member contributions, challenges, responsibilities, completed code, and deadlines.
Demonstration preparation	Prepare a timed code walkthrough and current-progress demo script with exact commands and evidence checkpoints.
Validation planning	Convert open technical gaps into measurable tests and blocking acceptance criteria.
Validation execution	Correct partition allocation and execute performance, concurrency, controlled-failure, startup, rollback, test, coverage, and dependency-audit checks.

6.2 Specific Work Completed by Daniel Grijalva

Completed activity	Detailed contribution
Project and document audit	Reviewed the Sprint 3 implementation document, Sprint 4 rubric, repository baseline, README, API orchestration, ML pipeline, speaker split logic, and open governance evidence.
Requirements update	Retained all 65 Sprint 2 requirements and documented twelve current clarifications without renumbering identifiers.
Risk update	Extended the register to fifteen risks and added evidence-freshness treatment for the gap between the July 8 repository baseline and July 20 submission.
Test-plan update	Expanded the detailed matrix to 32 tests, added a submission-freshness test, and preserved measurable blocking thresholds.
Implementation trace	Mapped learner UI, FastAPI, inference orchestration, storage adapters, ML pipeline, MLflow registry, Prometheus/Grafana, and CI to current work packages.
Professional deliverables	Prepared the comprehensive Markdown/PDF report, visible architecture/workflow diagrams, and a timed video walkthrough plan.
Validation hardening	Implemented the nonempty split allocator, added failure-injection regressions, and created a repeatable Python 3.11 validation harness.
Fresh evidence	Recorded 30-request prediction and health measurements, five-way concurrency, two clean starts, SHA-256 rollback, 45 passing tests, coverage, and dependency findings.

6.3 Challenges and Issues Addressed

Challenge	Treatment
Sole-member workload	All team roles are assigned to one person; work is prioritized around blocking acceptance gates and reproducible documentation.
Synthetic-only validation	The document separates functional fixture results from real-world accuracy claims and keeps real-data acceptance open.
Empty validation partition	Replaced floor-only allocation with deterministic minimum allocation and added a six-speaker regression test.

Challenge	Treatment
Slow end-to-end prediction	Replaced the single cold historical sample with 30 warm end-to-end measurements; p95 is 84.749 ms.
Open issue states	The report explicitly records the mismatch between evidence documents and unchecked/open GitHub issues.
Jira/Confluence wording	GitHub Issues and Project #4 are documented as the replacement; no unverified Jira or Confluence update is claimed.
Stale execution evidence	Refreshed executable, learner-UI, MLflow, Grafana, Prometheus, and GitHub Actions evidence; authenticated Project evidence remains open.
Vulnerable dependencies	Recorded all 32 fresh findings without silently accepting them; upgrade testing or signed risk acceptance remains blocking.

6.4 Responsibilities Assigned

Assignee	Role	Responsibilities
Daniel Grijalva	Owner and sole CLC member	Run the stack and tests; maintain code; update GitHub issues/project; collect evidence; resolve or accept risks; record and upload the demonstration; submit the PDF and hosted video URL.
Ryan Woodward	Instructor and evaluator	Review whether the GitHub replacement is acceptable, verify evidence, and evaluate unresolved blocking criteria.
Future maintainer	Deferred technical role	Add licensed real data, remediate dependencies, validate managed-service parity, and maintain the evidence automation.

6.5 Completed Coding Available for Demonstration

Code / configuration	Demonstrable responsibility
<code>api/app/main.py</code>	Creates the FastAPI application, mounts the learner UI, injects request IDs, records request metrics, and exposes <code>/metrics</code> .
<code>api/app/routers/predict.py</code>	Accepts <code>expected_phrase</code> and audio, records metadata-only logs, invokes inference, updates model/prediction metrics, and returns the typed response.
<code>api/app/services/inference.py</code>	Validates media, stores the upload temporarily, loads the approved model, normalizes audio/text, performs prediction, and deletes the upload.
<code>ml/pipeline/local_pipeline.py</code>	Runs <code>validate -> preprocess -> split -> features -> train -> evaluate -> gate -> model card -> MLflow register</code> .
<code>ml/data/split.py</code>	Performs deterministic speaker-aware splitting, guarantees requested nonempty partitions when possible, and computes a SHA-256 manifest hash.
<code>scripts/run_sprint4_validation.py</code>	Executes two startup cycles, 30 prediction and health requests, five-way concurrency, cleanup checks, and SHA-256 rollback.
<code>infrastructure/compose.yaml</code>	Defines the local service stack for API, storage, MLflow, Prometheus, and Grafana.
<code>.github/workflows/</code>	Runs format, lint, tests, secret scan, dependency audit, Compose validation, container smoke, and documentation checks.

6.6 Current and Future Deadlines

Date / window	Deadline	Expected output
July 20, 2026	Submit Sprint 4 PDF and hosted video URL	Current submission package
Before recording	Capture the current GitHub Project screen; record and upload the walkthrough	Owner-authenticated evidence checkpoint
Next implementation session	Acquire and document a licensed public validation subset; test a supported MLflow/pyarrow upgrade	Highest-priority technical work

Date / window	Deadline	Expected output
Before final release	Resolve or sign acceptance for dependency findings and reconcile GitHub Project issue fields	Remaining blocking acceptance work
August 9, 2026	Final reproducible release and presentation	Planned final milestone

6.7 Submission Readiness Statement

Ready for document submission: The report follows the required section structure, begins each major section on a new page, includes detailed implementation and validation plans, preserves the required Sprint 2 and Sprint 1 appendices, identifies the sole CLC member, and supports a separate hosted video demonstration.

Not yet final PoC acceptance: Local hardening, service screenshots, and GitHub Actions pass, but the document intentionally does not claim real-data validation, dependency acceptance, authenticated GitHub Project reconciliation, Confluence approval, or a hosted video. The video should demonstrate current progress and limitations rather than claim production readiness.

7. References

- Amazon Web Services. (n.d.-a). *Implement MLOps*. Amazon SageMaker AI Developer Guide. <https://docs.aws.amazon.com/sagemaker/latest/dg/mlops.html>
- Amazon Web Services. (n.d.-b). *What is the AWS Serverless Application Model (AWS SAM)?* <https://docs.aws.amazon.com/serverless-application-model/latest/developerguide/what-is-sam.html>
- DanielAndi. (2026). *FluentEdge* [Source code repository]. GitHub. <https://github.com/DanielAndi/fluentedge>
- Git. (n.d.). *git-rev-parse documentation*. <https://git-scm.com/docs/git-rev-parse>
- Git. (n.d.). *Hash function transition*. <https://git-scm.com/docs/hash-function-transition>
- GitHub. (n.d.). *About Projects*. GitHub Docs. <https://docs.github.com/en/issues/planning-and-tracking-with-projects/learning-about-projects/about-projects>
- Grijalva, D. (2026). *FluentEdge design and requirements specification* (Version 2.0) [Project document]. Grand Canyon University.
- Grijalva, D. (2026, June 28). *FluentEdge MLOps Sprint 1 proposal: AWS MLOps pipeline for AI-powered speaking feedback* [Project proposal]. Grand Canyon University.
- Grijalva, D. (2026, July 12). *FluentEdge prototype implementation document* (Version 3.3) [Project document]. Grand Canyon University.
- Mozilla Foundation. (n.d.). *Common Voice datasets*. <https://commonvoice.mozilla.org/en/datasets>
- National Institute of Standards and Technology. (2015). *Secure Hash Standard (SHS)* (FIPS PUB 180-4). <https://csrc.nist.gov/pubs/fips/180-4/upd1/final>

Appendix A - Sprint 2 Requirements Baseline

Source: The 65 requirement rows below are reproduced from
FluentEdge_Sprint2_Design_and_Requirements_Specification.md , version 2.0. Sprint 4 interpretations and changes appear in Section 2 so this appendix remains the baseline.

A.1 Data Input and Preprocessing Requirements

ID	Priority	Requirement	Acceptance criterion
FR-DI-001	Must	The system shall accept WAV, FLAC, MP3, or M4A input that can be decoded by the selected audio library.	Automated tests accept supported samples and reject unsupported or corrupt files.
FR-DI-002	Must	The API shall reject files larger than 10 MB and audio longer than 30 seconds for the prototype.	Boundary tests return documented 4xx responses.
FR-DI-003	Must	The system shall normalize training and inference audio to mono 16 kHz PCM before feature extraction.	Unit test verifies sample rate, channels, and output shape.
FR-DI-004	Must	The preprocessing stage shall validate required metadata fields, missing values, duplicate clip IDs, duplicate paths, and unreadable files.	Validation report identifies each failure category.
FR-DI-005	Must	Text normalization shall lowercase input, normalize whitespace, and apply documented punctuation rules.	Golden test cases match expected normalized text.
FR-DI-006	Must	Train, validation, and test splits shall be reproducible and speaker-aware when speaker identifiers are available.	Repeated runs with the same seed produce identical manifests and no speaker overlap.
FR-DI-007	Should	The pipeline should calculate transcript similarity, WER, CER, duration, speech rate, MFCC statistics, and signal-quality indicators.	Feature schema contains all implemented features with no invalid numeric values.

A.2 ML Model Requirements

ID	Priority	Requirement	Acceptance criterion
FR-ML-001	Must	The baseline model shall classify attempts as <code>pass</code> or <code>needs_review</code> .	Held-out predictions use only the defined labels.
FR-ML-002	Must	The training process shall log parameters, code version, dataset manifest hash, metrics, and artifact path.	MLflow run contains each required field.
FR-ML-003	Must	The project shall compare at least two baseline candidates, such as logistic regression and random forest or gradient boosting.	Evaluation report contains a side-by-side comparison.
FR-ML-004	Must	A candidate shall achieve macro F1 of at least 0.75 or produce a documented remediation decision.	Evaluation gate records pass or documented exception.
FR-ML-005	Must	The selected model shall report precision, recall, F1, confusion matrix, WER, CER, and inference latency.	Metrics file and report contain all required values.
FR-ML-006	Must	Model approval shall require human review and an explicit approved status in the local registry.	No model is marked production unless approval metadata exists.
FR-ML-007	Must	The inference service shall load a versioned artifact rather than an untracked local file.	Response includes model version and artifact checksum.
FR-ML-008	Should	The system should produce confidence or probability values and document calibration limitations.	Output schema includes confidence and report includes calibration note.
FR-ML-009	Should	Subgroup metrics should be reported only when metadata is present and the subgroup has at least 100 test examples.	Report suppresses undersized groups.

ID	Priority	Requirement	Acceptance criterion
FR-ML-010	Could	An embedding-based or small neural model may be evaluated as an extension.	Optional comparison is clearly labeled as experimental.

A.3 Local Infrastructure and AWS Portability Requirements

ID	Priority	Requirement	Acceptance criterion
FR-INF-001	Must	The complete local stack shall start with one documented command or a short documented command sequence.	A clean checkout can start all required services.
FR-INF-002	Must	Docker Compose shall define the API, storage service, MLflow, Prometheus, and Grafana services used by the prototype.	<code>docker compose config</code> succeeds and services pass health checks.
FR-INF-003	Must	The API shall support environment-based endpoints so local storage and future AWS S3 can use the same interface.	Tests run against a local endpoint without changing application code.
FR-INF-004	Must	AWS SAM CLI shall be used to test at least one Lambda/API Gateway-compatible request path locally or the document shall record why the FastAPI container was used instead.	<code>sam build</code> and <code>sam local start-api</code> evidence or approved design decision exists.
FR-INF-005	Must	The ML workflow shall run locally on a small dataset through a reproducible command.	Local pipeline produces cleaned data, model, and metrics artifacts.
FR-INF-006	Must	Local artifacts shall be stored in versioned directories or object keys that include dataset and run identifiers.	No approved artifact is overwritten by a later run.
FR-INF-007	Must	Infrastructure configuration shall contain no committed credentials.	Secret scan and repository inspection pass.
FR-INF-008	Must	The system shall provide health checks for the API, storage dependency, model availability, MLflow, Prometheus, and Grafana.	Health-check script exits successfully when the stack is ready.
FR-INF-009	Should	Infrastructure should be represented as code using Docker Compose plus AWS SAM, Terraform, or both.	Source-controlled infrastructure can recreate the local environment.
FR-INF-010	Should	The same inference image should be portable to an AWS container or SageMaker-compatible deployment.	Image exposes documented health and inference routes.
FR-INF-011	Should	GitHub Actions should validate Compose configuration and build the API image without requiring AWS credentials.	CI run completes on a pull request.
FR-INF-012	Could	An optional AWS deployment workflow may be added behind a manual approval and environment protection rule.	No cloud deployment occurs automatically from an unreviewed branch.

A.4 API, Audit, and Feedback Requirements

ID	Priority	Requirement	Acceptance criterion
FR-API-001	Must	<code>GET /health</code> shall return service status and dependency summaries without sensitive data.	Contract test validates status code and schema.
FR-API-002	Must	<code>POST /predict</code> shall accept an expected phrase and one supported audio file or object reference.	Valid request returns a documented response.

ID	Priority	Requirement	Acceptance criterion
FR-API-003	Must	The prediction response shall include score, label, confidence, feedback categories, model version, request ID, and latency.	OpenAPI schema requires all fields.
FR-API-004	Must	Invalid requests shall return deterministic 4xx errors with safe messages and request IDs.	Negative tests cover missing prompt, bad media, oversized file, and decode failure.
FR-API-005	Must	Raw audio, access tokens, and full private transcripts shall not appear in application logs.	Automated log-redaction test passes.
FR-API-006	Must	Each prediction shall emit aggregate metrics for request count, status, latency, and prediction label.	Prometheus endpoint exposes the documented metrics.
FR-API-007	Should	The API should generate interactive OpenAPI documentation in local development.	<code>/docs</code> opens and matches the committed contract.

A.5 GitHub Project and Collaboration Requirements

ID	Priority	Requirement	Acceptance criterion
FR-PM-001	Must	GitHub Issues shall represent sprint tasks with scope, acceptance criteria, owner, status, priority, requirement IDs, and evidence links.	All Sprint 2 issues contain the required fields or metadata.
FR-PM-002	Must	A GitHub Project shall replace Jira and contain all active and completed project issues.	Project link opens and contains the Sprint 2 issue set.
FR-PM-003	Must	The GitHub Project shall include Status, Priority, Sprint, Workstream, Requirement IDs, Target Date, and Evidence fields.	Field list or screenshot confirms each field.
FR-PM-004	Must	The project workflow shall use Backlog, Ready, In Progress, In Review, Blocked, and Done states.	Board screenshot shows the states or equivalent approved configuration.
FR-PM-005	Must	Issues and pull requests shall be linked to the corresponding requirements and evidence.	At least one end-to-end trace exists from requirement to issue to commit/test.
FR-PM-006	Must	GitHub Actions shall provide coding evidence for tests, linting, and build validation.	Workflow run links are present in the evidence index.
FR-PM-007	Should	Repository labels and milestones should identify workstream, priority, sprint, and risk.	Labels and milestones are visible in GitHub.
FR-PM-008	Should	Issue templates should require acceptance criteria and evidence before closure.	Template files exist and a sample issue uses them.
FR-PM-009	Should	The setup shall be reproducible through a documented GitHub CLI script.	Script is idempotent or safely detects existing resources.
FR-PM-010	Could	Automated Project field updates may be implemented with GitHub Actions or GraphQL after the baseline board is complete.	Automation is documented and does not expose tokens.

A.6 Non-Functional Requirements

ID	Category	Requirement	Measurement
NFR-PERF-001	Performance	Local <code>/predict</code> p95 latency shall be no more than 2.0 seconds for a 10-second prototype clip after warm-up.	At least 30 measured requests.
NFR-PERF-002	Performance	<code>GET /health</code> p95 latency shall be below 250 ms locally.	Automated performance test.
NFR-REL-001	Reliability	The API shall return a controlled error instead of crashing when a dependency is unavailable.	Forced dependency-failure test.
NFR-REL-002	Reliability	The last approved model artifact shall remain available as the rollback target.	Rollback test loads the previous version.

ID	Category	Requirement	Measurement
NFR-REL-003	Reliability	Local stack startup shall succeed on two consecutive clean runs.	Startup logs and health-check output.
NFR-SCALE-001	Scalability	The API shall support at least five concurrent local requests without data corruption.	Concurrency test passes.
NFR-SEC-001	Security	Secrets shall be loaded from environment variables or approved secret stores and never committed.	Secret scanner and repository review.
NFR-SEC-002	Security	Uploaded filenames shall be replaced with generated object keys and shall not control filesystem paths.	Path-traversal test passes.
NFR-SEC-003	Security	Dependencies shall be pinned or bounded and scanned for known vulnerabilities.	Dependency scan attached.
NFR-PRIV-001	Privacy	Raw uploaded audio shall be deleted within 24 hours by default and sooner after successful inference when configured.	Cleanup test and configuration evidence.
NFR-PRIV-002	Privacy	Logs shall exclude raw audio, access tokens, and full user transcripts.	Log-redaction test.
NFR-OBS-001	Observability	Metrics shall include request count, error count, p50/p95 latency, prediction distribution, confidence summary, and active model version.	Dashboard and metric endpoint evidence.
NFR-OBS-002	Observability	Every request shall have a correlation/request ID.	API and log sample.
NFR-PORT-001	Portability	The API and inference service shall run from the same Docker image in local and AWS-targeted configurations.	Local container smoke test.
NFR-MAINT-001	Maintainability	Source code shall pass formatting, linting, type checks where configured, and unit tests in CI.	GitHub Actions run.
NFR-MAINT-002	Maintainability	Public functions and configuration files shall include clear documentation and examples.	Code review checklist.
NFR-USAB-001	Usability	Learner feedback shall use plain language and shall not claim clinical pronunciation accuracy.	UI review against approved wording.
NFR-COST-001	Cost	The required local demonstration shall not require paid AWS resources.	Local demo completes with no cloud resources.
NFR-COST-002	Cost	Any optional AWS resources shall have documented shutdown and deletion steps.	Operations checklist.

Appendix B - Sprint 1 Risk Baseline

Source: The eight rows below reproduce the risk assessment from page 13 of `mlops_sprint1_proposal_daniel_grijalva-1.pdf`, dated June 28, 2026. The B-R identifiers are added only for cross-referencing in this implementation document; the risk wording, qualitative impact/probability, mitigation, and contingency are preserved from Sprint 1.

ID	Sprint 1 risk	Impact	Probability	Sprint 1 mitigation	Sprint 1 contingency plan
B-R1	Dataset labels do not perfectly represent pronunciation quality.	High	Medium	Use a <code>pass / needs-review</code> prototype label, document assumptions, and include WER/CER and acoustic features instead of claiming expert pronunciation grading.	Reduce scope to scoring expected phrase similarity and confidence.
B-R2	Accent or locale underperformance creates unfair feedback.	High	Medium	Evaluate subgroup performance where metadata exists; report metadata limitations; avoid overclaiming fairness.	Add threshold tuning or route low-confidence cases to <code>needs-review</code> .
B-R3	Model accuracy does not meet target.	Medium	Medium	Start with interpretable baseline, inspect confusion matrix, improve features incrementally.	Lower target with documented reason or compare multiple models.
B-R4	AWS costs exceed student budget.	High	Medium	Use small sample data, local development first, budgets/alarms, and shutdown instructions.	Use local Docker demo and AWS architecture evidence if endpoint is too expensive.
B-R5	Deployment endpoint has high latency.	Medium	Medium	Optimize audio preprocessing, use small model, test p95 latency, and reduce input size.	Deploy batch/demo endpoint instead of real-time endpoint.
B-R6	Data drift appears after deployment.	High	Medium	Monitor prediction distribution, feature drift, WER/CER trends, and confidence shifts.	Open retraining ticket and roll back to previous approved model if necessary.
B-R7	Scope creep from trying to build a full app.	Medium	High	Keep scope limited to MLOps prototype and REST API.	Move UI/mobile features to future work.
B-R8	Team/collaboration evidence is incomplete.	High	Low	Maintain Jira, Confluence, GitHub issues/commits, review notes, and screenshots each sprint.	Create a Sprint Evidence Appendix before final submission.

Sprint 1 risk threshold: Any risk marked High impact must have either a mitigation action, a contingency plan, or both before implementation starts.

B.1 Cross-Reference to Current Assessment

Baseline	Current risk(s)	Change
B-R1	R-01, R-02	Remains Critical because synthetic-only data and perfect fixture scores prevent real-world validity claims, despite the corrected nonempty split.
B-R2	R-03	Remains Critical until a sufficiently large, licensed, metadata-bearing dataset supports subgroup evaluation.
B-R3	R-01, R-02	The fixture gate passes, but real-data accuracy and leakage-resistant validation remain unresolved.
B-R4	R-06, R-13	Reduced because the required implementation is local-first; parity and optional cloud-cost controls remain.
B-R5	R-04	Reduced after the 30-request warm run measured p95 84.749 ms; larger and real-data traffic remains untested.
B-R6	R-07, R-14	Rollback risk is reduced by the verified restore exercise; production drift evidence remains unavailable.
B-R7	R-09	Controlled through a limited REST prototype, deferral of optional neural/AWS work, and prioritized acceptance gates.

Baseline	Current risk(s)	Change
B-R8	R-10, R-15	Remains High/Critical because issue states, evidence age, and current submission claims must be reconciled.

B.2 Risks Added After Sprint 1

Current risk	Reason it was added
R-05 Dependency vulnerabilities	Sprint 2 dependency scanning identified actionable MLflow/pyarrow advisories.
R-08 Privacy or retention failure	Implementation introduced temporary audio storage, generated object keys, log-redaction, and deletion controls that require explicit risk tracking.
R-11 Reproducibility hash ambiguity	Instructor feedback required explicit SHA versions and full authoritative digests.
R-12 Service availability	The multi-service local stack introduced dependency health, controlled-failure, and repeatable-startup acceptance concerns.

Appendix C - Current Evidence Index

Purpose: Identify the repository and submission evidence available for the Sprint 4 demonstration and the exact use of each artifact.

Evidence area	Artifact / result	Demonstration use
Repository baseline	Sprint 4 implementation 5c191a638110e28ceff15e259766bc1c04dc9d99 ; green Docs fix e68e8f0d766701f76747b221ebd70f2ffeb2ab62	Run <code>git rev-parse HEAD</code> at the start of the video.
GitHub Actions	CI, Container, and Docs runs completed successfully	Show <code>actions-ci-sprint4.png</code> and the linked workflow runs.
Quick-start documentation	README.md: <code>cp .env.example .env</code> ; <code>make setup</code> ; <code>make up</code> ; <code>make bootstrap</code> ; <code>make health</code>	Show the README and terminal health output.
API application	<code>api/app/main.py</code>	Show request-ID middleware, metrics endpoint, and router registration.
Prediction route	<code>api/app/routers/predict.py</code>	Show form/file input, metadata-only logging, metrics, and response fields.
Inference orchestration	<code>api/app/services/inference.py</code>	Show media validation, temporary storage, approved model load, normalization, prediction, and cleanup.
ML pipeline	<code>ml/pipeline/local_pipeline.py</code>	Show validation, preprocessing, split, features, two candidates, quality gate, model card, and registration.
Corrected split behavior	<code>ml/data/split.py</code> , <code>tests/unit/test_split.py</code>	Show deterministic 4/1/1 speaker-group allocation and the impossible-partition guard.
Fresh automated evidence	45 tests passed, 2 skipped; Ruff format/lint pass; 78% coverage	Show <code>docs/evidence/E-11-sprint4-validation/</code> .
Fresh performance evidence	<code>/health</code> p95 36.732 ms; warm <code>/predict</code> p95 84.749 ms over 30 requests	Show <code>validation-report.json</code> and startup logs.
Reliability evidence	Five simultaneous requests, controlled storage/model 503 tests, two startup cycles, SHA-256 rollback	Show the validation report and JUnit output.
Observability	Fresh Prometheus, Grafana, MLflow, and learner-UI screenshots	Show the healthy target, populated dashboard, registered runs, and successful prediction.
Submission artifacts	Sprint 4 Markdown/PDF and hosted video walkthrough	Upload the video to a platform and submit its URL rather than the video file.

Before upload: Replace historical screenshots or statements with fresh evidence where practical, verify that the video link is accessible without requesting permission, and submit the PDF plus hosted URL.